

# ROHAN DHANRAJ YADAV

AI/ML Engineer | Generative AI | Multi-Agent Systems | AI Platform Developer

- ✉ Email: [rohan.dhanraj.y@gmail.com](mailto:rohan.dhanraj.y@gmail.com)  
☎ Phone: [+91-7008958143](tel:+91-7008958143)  
🌐 LinkedIn: [linkedin.com/in/rohan-dhanraj-yadav](https://linkedin.com/in/rohan-dhanraj-yadav)  
🐙 GitHub: [github.com/rohandhanraj](https://github.com/rohandhanraj)  
🌐 Portfolio: [rohandhanraj.github.io](https://rohandhanraj.github.io)  
📍 Location: Bengaluru, India

## PROFESSIONAL PROFILE

Results-driven AI/ML Engineer with **7+ years of professional experience** including **4+ years specializing in designing, developing, and deploying production-grade AI systems** across enterprise domains. Deep expertise in **Generative AI, Agentic AI, Large Language Models (LLMs), Retrieval-Augmented Generation (RAG), Multi-Agent Orchestration**, and **Machine Learning**.

Proven track record in leading cross-functional teams of **15+ engineers**, architecting scalable ML/AI pipelines, and delivering high-impact AI applications for **Fortune 500 clients** including Cisco, ITC Fairview, KDP, Omnitrax, PDS, and Marco Technologies. Known for building production systems handling **4,500+ daily API requests**, processing **100K+ daily records**, and achieving **18-38% model performance improvements** in production environments.

### Career Metrics

| Total Professional Experience | AI/ML Engineering Experience | Agentic AI Experience | Generative AI Experience | Production AI Projects Delivered | Max Team Size Led | Fortune 500 Clients Served |
|-------------------------------|------------------------------|-----------------------|--------------------------|----------------------------------|-------------------|----------------------------|
| 7+ years                      | 4+ years                     | 1.5 year              | 3 year                   | 8+                               | 15 engineers      | 6+                         |

## TECHNICAL SKILLS

### Programming Languages & Frameworks

Python, Django, Flask, FastAPI, Streamlit

### Generative AI & LLM Engineering

**LLM Providers & Platforms:** OpenAI, Anthropic, Vertex AI, Google Gen AI Platform, Cohere, AI21 Labs

**LLM Frameworks:** LangChain

**Prompt Engineering & Reasoning:** Prompt Engineering, Prompt Chaining, Chain-of-Thought (CoT) Reasoning, Structured Output Generation

**Retrieval & Context Engineering:** Retrieval-Augmented Generation (RAG), Embedding Generation

**LLM Architecture & Solution Design:** LLM System Design, LLM Solution Architecture

### Agentic AI & Multi-Agent Systems

**Agent Frameworks:** LangGraph

**Agent Design Patterns:** ReAct, Reflexion, Self-Ask, Plan-and-Execute, Tool Use, Human-in-the-Loop, Tree of Thoughts (ToT), Multi-Agent Collaboration, Deep Research

**Agent Development & Orchestration:** Agent Orchestration, LangGraph Workflows, Multi-Agent Systems, Tool Integration, Memory Management, Autonomous AI Agents, Autonomous Decision-Making

## Machine Learning & Deep Learning

**Machine Learning Libraries:** Scikit-Learn, PyCaret

**Deep Learning Frameworks:** TensorFlow, Keras, PyTorch

**Model Development & Optimization:** Neural Networks, Hyperparameter Optimization, Optuna

## Model Fine-Tuning & Alignment

**Parameter-Efficient Fine-Tuning (PEFT):** PEFT, LoRA, QLoRA, IA<sup>3</sup>, LLaMA-Adapter

**Fine-Tuning Methodologies:** Instruction Fine-Tuning, Supervised Fine-Tuning (SFT), Soft Prompting

**Alignment & Preference Optimization:** Reinforcement Learning from Human Feedback (RLHF), Reinforcement Learning from AI Feedback (RLAIF), Direct Preference Optimization (DPO)

## Natural Language Processing (NLP)

**NLP Libraries:** SpaCy, NLTK, Hugging Face Transformers

**Encoder-Based Transformer Applications:** BERT, RoBERTa, Sentiment Analysis, Aspect-Based Sentiment Analysis (ABSA), Named Entity Recognition (NER), Text Classification

**Decoder-Based Transformer Applications:** Zero-Shot Learning, One-Shot Learning, Few-Shot Learning, Text Generation, Summarization, Question Answering, Paraphrasing

**Encoder-Decoder Transformer Applications:** Machine Translation

## Data Engineering & Databases

**Relational Databases:** MySQL, PostgreSQL, BigQuery

**NoSQL Databases:** MongoDB, Cassandra

**Vector Databases:** Pinecone, Qdrant, LanceDB, Chroma, PgVector, Milvus, Weaviate, Sqlitevec

**Graph Databases:** Neo4j

**Data Engineering:** Apache Airflow, ETL Pipelines, Kafka

## Analytics & Statistical Computing

NumPy, Pandas, SciPy, Matplotlib, Seaborn, Plotly, Statistical Analysis, Predictive Analytics, Data Visualization

## Web Scraping & Data Extraction

**Web Scraping Frameworks:** BeautifulSoup, Scrapy, Selenium, Playwright

**Data Extraction & Validation:** Data Extraction, Data Validation

## DevOps, Cloud & MLOps

**Version Control & DevOps:** Git, GitHub, Docker

**CI/CD & Automation:** GitHub Actions, Jenkins

**Cloud Platforms:** Google Cloud Platform (GCP), Amazon Web Services (AWS)

**MLOps Tools:** MLflow, TensorBoard, DVC

**MLOps Practices:** Continuous Model Training, Model Evaluation, Model Deployment, Model Monitoring, Model Optimization, Model Versioning, Model Management

---

## WORK EXPERIENCE

### 1. Senior AI/ML Engineer — Generative AI | Agentic AI | MCP Tools

**R Systems International** | June 2026 – Present | Full-time

**Client:** Lendistry

**Focus:** Enterprise Agentic AI workflows, MCP tools integration, guardrails, harness engineering, observability, and Agentic Ops

#### Key Responsibilities & Achievements:

- Architected enterprise-grade **Agentic AI workflows** and multi-agent systems using Model Context Protocol (MCP) and custom guardrails for fintech operations
- Implemented robust **Harness Engineering** and **Agentic Ops** pipelines for automated tool execution, schema validation, and real-time observability
- Developed **ML Flow** integration and **Agentic E2E evaluation** frameworks for continuous testing, strategic agent evaluation, and regression tracking
- Pioneered safety guardrails and multi-layered retrieval backstops to ensure reliable autonomous decisions in production workflows

**Tech Stack:** Python, Generative AI, Agentic Workflows, MCP Tools, Guardrails, Harness Engineering, Agentic Ops, ML Flow, Agentic E2E Evaluation, Observability

---

### 2. AI/ML Engineer — Generative AI | Multi-Agent Systems

**AI Tech Solutions Ltd.** | November 2024 – October 2025 | Full-time

**Focus:** Architecting production-ready Agentic AI platforms, advanced RAG systems, and multi-agent orchestration

#### Key Responsibilities & Achievements:

- Led 15-member team building scalable AI systems using Python, FastAPI for backend services and production deployment
- Developed modular ML pipelines (Omodore platform) achieving **38% accuracy improvement** through advanced RAG architecture
- Designed multi-agent orchestration systems using **LangGraph** for complex workflow automation and autonomous decision-making
- Implemented prompt chaining and **chain-of-thought reasoning** mechanisms
- Deployed containerized AI systems using **Docker** with vector databases (Qdrant, LanceDB, Chroma)
- Built automated web scraping systems (Rbetraj) with **Playwright** and **Selenium** for real-time data extraction
- Architected multi-modal virtual agent (GALAMBO) integrating GPT-4 and autonomous tool use

#### Projects

- **Omodore** — Autonomous Agentic AI platform demonstrating LangGraph state orchestration
- **GALAMBO** — Multi-modal AI assistant integrating vision models and real-time tool use
- **Rbetraj** — Automated Dropshipping Platform with AI-Powered Product Research and Listing Generation by scanning Global E-Commerce Sites.

**Tech Stack:** Python, LangChain, LangGraph, FastAPI, LLMs, RAG, Agent Orchestration, Multi-Agent Systems, Docker, Qdrant, LanceDB, Chroma, Playwright, Selenium

---

### 3. Software Engineer — Generative AI | ML | Data Science | Python

Aimlytics Technology | September 2022 – October 2024 | Full-time

**Focus:** Enterprise-scale Generative AI, NLP, and ML solutions for Fortune 500 clients

#### Key Responsibilities & Achievements:

- Led team of 8 engineers delivering AI solutions for **Fortune 500 clients**
- Built RAG-based virtual sales assistant (SalesMoji) handling **4,500+ daily API requests**
- Built ML pipelines processing **100K+ daily records** using **Apache Airflow**
- Fine-tuned **BERT** and **RoBERTa** models achieving **18% performance improvement**
- Implemented vector database solutions (Qdrant, Chroma) for semantic search
- Built multi-modal educational AI system (InventtEd)
- Developed multilingual NLP pipeline supporting **15+ languages** (evalmyBRAND)
- Resolved **28 critical production issues** with 100% on-time delivery

#### Projects

- **SalesMoji** — AI-Powered Lead Generation Platform with Context-Aware Follow-up Automation and Real-time LinkedIn Integration.
- **InventtEd** — Multi-Modal AI Learning Assistant Enhancing STEM Education with Visual Explanations and Interactive Problem-Solving
- **evalmyBRAND** — Brand Monitoring and Reputation Management Platform with Multi-Lingual Sentiment Analysis and Trend Tracking

**Tech Stack:** Python, RAG, LangChain, Airflow, BERT, RoBERTa, FastAPI, Flask, TensorFlow, PyTorch, Qdrant, Chroma, GCP

---

### 4. Machine Learning Engineer Intern

iNeuron.ai | August 2021 – September 2022 | Internship

**Focus:** ML pipeline development and optimization

#### Key Responsibilities & Achievements:

- Built ML pipeline architectures for multi-class classification
- Reduced training time by **35%** using **Optuna**
- Deployed ML applications using Flask and Docker
- Built protein classification system (77 markers, 8 classes)

#### Projects

- **Protein Classification System** — Multi-Class Classification Model for High-Throughput Protein Analysis using ML Pipeline Architectures.

**Tech Stack:** Python, Flask, Scikit-Learn, Optuna, MySQL, Cassandra, Docker, AWS

---

### 4. Metallurgical Quality Engineer

Bureau Veritas Group | August 2016 – November 2019 | Full-time

**Focus:** Quality assurance and industrial data analytics

#### Key Responsibilities & Achievements:

- Conducted metallurgical testing and quality assurance across major steel companies
  - Performed statistical analysis, predictive analysis, and data-driven quality improvement
  - Applied Python (NumPy, Pandas, SciPy, Plotly) for industrial data analysis
  - Managed quality compliance across multiple facilities
- 

## PROJECT SYNOPSIS

### 1. Project Name: OMODORE

**Duration:** November 2024 to October 2025

**Role:** Senior AI Engineer

**Description:**

**Domain:** AI Agent Platform

Omodore is an auto-agent platform enabling non-technical users to build domain-specific AI agents without coding expertise. The platform leverages advanced RAG pipelines to deliver accurate, context-aware responses with a 45% accuracy boost.

**Key Features:**

- No-code AI agent builder for domain-specific applications
- Advanced RAG pipeline with 45% accuracy improvement
- Context-aware intelligent responses using Chain of Thoughts
- Scalable architecture for enterprise deployment
- Automated agent workflows

**Roles and Responsibilities:**

- Led 15-member team developing Gen AI and Agentic AI solutions
- Architected scalable RAG pipelines using LangChain and LangGraph
- Implemented prompt chaining and Chain of Thoughts methodology for improved reasoning
- Built automated agent building workflows using FastAPI
- Integrated multiple LLMs (Large Language Models) for optimal performance
- Implemented vector database solutions (Qdrant, Milvus, Weaviate) for efficient information retrieval
- Resolved critical production issues ensuring system reliability

**Skills Used:**

**Language:** Python

**Development Framework:** FastAPI

**Machine Learning:** Sklearn, TensorFlow, Keras, PyTorch

**NLP:** Hugging Face Transformers, Large Language Models (LLMs)

**Model Development:** Model Tuning, Model Training, Model Evaluation, Model Deployment, Model Monitoring, Model Optimization, Model Versioning, Model Management

**Generative AI & Agentic AI:** LangChain, LangGraph, GPT, LLMs, Prompt Chaining, Chain of Thoughts, RAG, Agent Building, Autonomous Agents, Multi-Agent Systems

**Data Engineering & Databases:** Apache Airflow, Qdrant, MongoDB, PostgreSQL, PgVector, Neo4j

**Cloud & DevOps:** Docker

---

### 2. Project Name: RBETRAJ

**Duration:** November 2024 to October 2025

**Role:** Senior AI Engineer

**Description:**

**Domain:** E-commerce Data Intelligence

Rbetraj is an Automated Dropshipping Platform powered by Shop Scan Agent which specializes in e-commerce data extraction with advanced price tracking and competitive analysis capabilities. At the heart of Shop Scan Agent is an Open Source Automated Scraper AI TechCrawler that intelligently scrapes websites, extracts structured data, and provides actionable insights.

**Key Features:**

- Automated website scraping with intelligent data extraction
- Structured output generation from unstructured web data using Pydantic
- Real-time price tracking across multiple e-commerce platforms
- Competitive analysis and market intelligence
- Multi-source data aggregation

**Roles and Responsibilities:**

- Developed open-source automated scraping framework [AI Tech Crawler](#)
- Built Shop Scan Agent for e-commerce intelligence
- Implemented GPT-powered and LLM-based structured data extraction
- Created robust scraping pipelines using Playwright
- Designed RESTful APIs using FastAPI for data access
- Integrated LangChain for intelligent data processing workflows
- Implemented Pydantic models for data validation and structured output
- Implemented error handling and retry mechanisms for reliable scraping

**Skills Used:**

**Language:** Python

**Development Framework:** FastAPI

**Web Scraping:** Selenium, Playwright, BeautifulSoup

**Generative AI & Agentic AI:** GPT, LLMs, LangChain, Pydantic for structured output **Data Engineering &**

**Databases:** Apache Airflow, Qdrant, MongoDB

**Cloud & DevOps:** Docker

---

### 3. Project Name: GALAMBO

**Duration:** November 2024 to October 2025

**Role:** Senior AI Engineer

**Description:**

**Domain:** Multi-Modal AI Assistant

GALAMBO is a sophisticated multi-modal virtual agent that combines context-aware GPT responses with advanced image search capabilities, providing users with comprehensive visual and textual information retrieval.

**Key Features:**

- Multi-modal AI processing (text and images)

- Context-aware conversational responses using LLMs
- Intelligent image search and analysis
- Integration with search engine tools
- Real-time information retrieval

**Roles and Responsibilities:**

- Architected multi-modal AI system integrating text and image processing
- Implemented context-aware response generation using GPT and LLMs
- Integrated search engine tools for enhanced information retrieval
- Built scalable API infrastructure using FastAPI
- Implemented LangChain workflows for agent orchestration
- Optimized response times and accuracy for production use

**Skills Used:**

**Language:** Python

**Development Framework:** FastAPI

**Generative AI & Agentic AI:** GPT, LLMs, Multi-modal AI, LangChain, Agent Building

**Tools:** Search Engine Integration

**Cloud & DevOps:** Docker

---

#### 4. Project Name: SALESMOJI

**Duration:** November 2023 to October 2024

**Role:** Software Engineer (Gen AI | ML | Data Science | PYTHON)

**Description:**

**Domain:** Sales and Marketing

SalesMoji harnesses the power of GPT and Large Language Models to understand sales data, answer questions in plain English, and deliver instant, actionable sales intelligence. The platform processes 1M+ daily records through sophisticated Airflow pipelines, transforming raw data into strategic insights.

**How it Works:** Chat-bot interface where users can type sample questions and receive instant data analytics with visual representations and actionable recommendations.

**Key Features:**

- AI-Powered Sales Insights using GPT and LLMs
- User-Friendly Conversational Interface
- Customizable Templates and Dashboards
- Real-Time Data Visualization
- Interactive Chat Interaction with 10K+ daily queries
- 89% accuracy in query responses
- 1M+ daily records processing capability

**Roles and Responsibilities:**

- Led 8-member team delivering custom AI workflows for Fortune 500 clients
- Built RAG-based Virtual Sales Assistant responding to natural language queries
- Achieved 10K+ daily query handling with 89% accuracy using context-aware GPT responses
- Developed RAG pipelines for various sales channels using Vector Databases
- Constructed and orchestrated data pipelines processing 1M+ daily records using Airflow
- Architected Machine Learning pipelines for predictive analytics

- Generated predictive insights for sales forecasting and trend analysis
- Implemented vector database solutions (Qdrant, Milvus, Weaviate) for efficient retrieval
- Applied Prompt Chaining and Chain of Thoughts for complex reasoning
- Built agent workflows using LangChain and LangGraph
- Resolved critical production issues maintaining high system availability

#### **Skills Used:**

**Language:** Python

**Development Framework:** FastAPI

**Machine Learning:** Sklearn, TensorFlow, Keras, PyTorch

**NLP:** Hugging Face Transformers, Large Language Models (LLMs)

**Generative AI & Agentic AI:** LangChain, LangGraph, GPT, Prompt Chaining, Chain of Thoughts, RAG, Agent Building

**Data Engineering & Databases:** Apache Airflow, MySQL, MongoDB, Qdrant, Milvus, Weaviate

**Cloud & DevOps:** Docker

## **5. Project Name: AISERA**

**Duration:** July 2023 to October 2024

**Role:** Software Engineer (AI | ML | Data Science | PYTHON)

#### **Description:**

**Domain:** Customer Service and IT Operations Management

Aisera is a leading provider of Generative AI Solutions that helps enterprises boost revenue, improve user productivity, lower operating expenses, and create magical user experiences by automating and personalizing customer support. The platform uses Natural Language Processing (NLP) and Machine Learning (ML) to understand customer queries and provide relevant information or solutions.

#### **Key Features:**

- Custom AI models for various enterprise tasks
- AI Service Management (AISM) for automated operations support
- TRAPS Framework for ethical and secure AI practices
- Positive customer feedback on efficiency and automation
- Enterprise-grade workflows for Fortune 500 clients

#### **Roles and Responsibilities:**

- Led 8-member team comprising 5 Data Engineers and 3 Software Developers
- Delivered custom AI workflows for Fortune 500 clients including Cisco, ITC Fairview, Omnitrax, PDS, KDP, Marco Technologies
- Fine-tuned BERT models for sentiment analysis achieving 50% performance improvement
- Developed and scheduled system triggers using Airflow for automated operations
- Implemented custom field prediction using advanced BERT architectures
- Architected scalable ML pipelines ensuring reliability and performance
- Implemented vector database solutions (Chroma, Weaviate) for RAG applications
- Integrated LangChain for building conversational AI workflows
- Resolved critical production issues maintaining high system availability
- Handled 28 tickets overall and individually resolved 16 tickets out of 63 tickets assigned to Aimlytics Technologies Workflow team

#### **Skills Used:**



**Language:** Python, Nashorn JAVA

**Development Framework:** FastAPI

**Machine Learning:** Sklearn, TensorFlow, Keras, PyTorch

**NLP:** Spacy, NLTK, Hugging Face Transformers, BERT, Large Language Models (LLMs)

**Generative AI & Agentic AI:** LangChain, Prompt Chaining, RAG, Agent Building, Autonomous Agents, Multi-Agent Systems, Agentic Guardrails and Middleware for Enterprise AI

**Data Engineering & Databases:** Apache Airflow, MySQL, MongoDB, Chroma, Weaviate

**Cloud Platform:** GCP

---

## 6. Project Name: EVALMYBRAND

**Duration:** September 2022 to October 2024

**Role:** Software Developer (AI | ML | Data Science | PYTHON)

### Description:

**Domain:** Customer Experience Management

evalmyBRAND is an AI-powered platform designed to turn unstructured customer feedback data into actionable insights, providing a holistic view of customer sentiment and brand perception across multiple channels. The platform supports 15+ languages for comprehensive global brand monitoring.

### Key Features:

- Comprehensive Insights and Analytics
- Customer Engagement Tracking
- Aspect-Based Sentiment Analysis with 50% performance gain
- Competitor Analysis and Benchmarking
- Brand Goals Attainment Monitoring
- Intelligent Case Management System
- Multi-language support (15+ languages)

### Roles and Responsibilities:

- Developed web application using FastAPI for scalable API services
- Built and orchestrated data pipelines using Airflow for automated data processing
- Managed team of 5 Data Annotators for dataset labeling and quality control
- Created innovative pipeline for mixed code language translation using Google Translate API and LibreOffice
- Fine-tuned RoBERTa model for Aspect-Based Sentiment Analysis achieving 50% performance improvement
- Implemented BERT models for Neural Machine Translation tasks with 27% accuracy improvement
- Optimized model architectures for production deployment
- Integrated multiple data sources for comprehensive brand analysis across 15+ languages
- Implemented sentiment analysis pipelines for real-time brand monitoring

### Skills Used:

**Language:** Python

**Development Framework:** Poetry, Flask, FastAPI

**Machine Learning:** Sklearn, TensorFlow, Keras, PyTorch

**NLP:** Spacy, NLTK, Hugging Face Transformers, BERT, RoBERTa, Large Language Models (LLMs), Google Translate, LibreOffice Translate, Mixed Code Language Translation

**Generative AI:** LangChain, Prompt Chaining, RAG

**Data Engineering & Databases:** Apache Airflow, Oracle, MongoDB, Cassandra

---

## 7. Project Name: INVENTED

**Duration:** January 2023 to August 2023

**Role:** Software Developer (AI | ML | Data Science | PYTHON)

### Description:

**Domain:** Educational Technology

InventtEd™ is a pioneering education platform designed to help educators and students become more efficient, effective, and productive. It provides infinite ways to improve productivity, enhance student learning outcomes, and increase engagement through AI-powered tools.

### Key Features:

- Intelligent Text Summarization
- Automated Essay Grading with detailed feedback
- Feedback/Complaint Analysis for institutional improvement
- Unusual Language Detection for content moderation
- Automated Timetable Generator
- Virtual Tutor (Q&A Bot) for student support
- Student Performance Prediction analytics

### Roles and Responsibilities:

- Performed advanced prompt engineering to fetch questions and answers from ChatGPT using optimized prompts
- Fine-tuned CLIP (Multi-Modal Vision Transformer) to generate images from text descriptions
- Fine-tuned Flan-T5 model to generate keyword-based MCQs from given educational context
- Designed scalable Phishing Classifier using classical Machine Learning algorithms for email security
- Built multi-modal model for social media content classification (Input: post text + image, Output: category)
- Integrated multiple AI models into cohesive educational platform
- Optimized Transformer models for real-time educational applications

### Skills Used:

**Language:** Python

**Development Framework:** Poetry, Flask, FastAPI

### Databases:

- SQL: Oracle
- NoSQL: MongoDB, Cassandra

**Machine Learning:** Sklearn, TensorFlow, Keras, PyTorch

**NLP:** Spacy, NLTK, Hugging Face Transformers, Large Language Models (LLMs), CLIP, Flan-T5

**Generative AI:** LangChain, Prompt Chaining, RAG, Agent Building, Prompt Engineering

**Translation Tools:** Google Translate, LibreOffice Translate, Mixed Code Language Translation

**Cloud Platform:** GCP

---

## 8. Project Name: MICE PROTEIN EXPRESSION

**Duration:** August 2021 to September 2022

**Role:** Machine Learning Engineer Intern

**Description:**

**Domain:** Bioinformatics and Protein Classification

Protein Expression classification models are frequently viewed not only as a challenging task but also as a classification problem that requires balancing accuracy and efficiency in analysis validation due to large-scale datasets.

This project analyzed expression levels of 77 proteins measured in the cerebral cortex of 8 classes of control and Down syndrome mice exposed to context fear conditioning—a task used to assess associative learning. The primary aim was to identify subsets of proteins that are discriminant between the classes, making this a complex multi-class classification problem.

**Key Features:**

- Multi-class classification across 8 distinct classes
- Analysis of 77 protein expression markers
- Automated ML pipeline for model selection
- Hyperparameter optimization for optimal performance
- 40% training time reduction via Optuna

**Roles and Responsibilities:**

- Led team of 10 members in developing ML solution
- Projected and architected comprehensive Machine Learning System
- Built automated ML architectures for efficient model development
- Automated Machine Learning pipelines for hyperparameter tuning achieving 40% training time reduction via Optuna
- Implemented model selection algorithms to identify best performing models for production deployment
- Developed data preprocessing pipelines for protein expression data
- Created scalable architecture for handling large biological datasets
- Deployed scalable ML applications with Flask web framework
- Integrated Cassandra DB and MySQL for efficient data management

**Skills Used:**

**Language:** Python

**Development Framework:** Flask

**Databases:** Cassandra DB, MySQL

**Machine Learning:** Sklearn (Scikit-Learn), Optuna for hyperparameter tuning, PyCaret **Tools:** Anaconda, Data Pipeline, Artificial Intelligence, Machine Learning

**Cloud & DevOps:** AWS, Docker

---

**EDUCATION**

| Degree                         | Institution                                   | Year |
|--------------------------------|-----------------------------------------------|------|
| B.Tech. Mechanical Engineering | Biju Patnaik University of Technology, Odisha | 2016 |

**GATE Qualified** — Credential ID: A2WR3

---

## CERTIFICATIONS

| Certification                    | Issuer        | Credential ID       |
|----------------------------------|---------------|---------------------|
| GATE ME 2016                     | IIT/NTA       | A2WR3               |
| Machine Learning Zoom Camp       | DataTalksClub | 8F163B              |
| Gold Badge: Python               | HackerRank    | 6920934A49FA        |
| Gold Badge: Problem Solving      | HackerRank    | 113775E95540        |
| Gold Badge: Statistics (10 Days) | HackerRank    | —                   |
| 30 Days of ML Challenge          | Kaggle        | Rank 69/7,573 teams |

## ACHIEVEMENTS & AWARDS

- **7+ years** of professional experience including 4+ years in AI/ML engineering
- Achieved **Gold Badges** for Python, Problem Solving, and Statistics on HackerRank
- Secured **Rank 69 out of 7,573 teams** in Kaggle ML Challenge
- Led teams of **8-15 members** across enterprise AI projects
- Delivered **50% model performance improvements** in production systems
- Built systems handling **4,500+ daily API requests**
- Processed **100K+ daily records** via Airflow pipelines
- Achieved **38% accuracy boost** in RAG systems
- Reduced training time by **35%** using Optuna
- Resolved **28 critical production issues**

Generated: April 2026 Location: Bengaluru, India