

ROHAN DHANRAJ YADAV

AI/ML Engineer | Generative AI | Agentic AI Systems

✉ rohan.dhanraj.y@gmail.com | ☎ +91-7008958143 | 📍 Bengaluru, India

🌐 LinkedIn | 📁 Portfolio | 🐙 GitHub

SUMMARY

Results-driven AI/ML Engineer with **7+ years of total professional experience**, including **4+ years** building production-grade AI systems. Expert in **Generative AI, Agentic AI, LLMs, RAG**, and **Multi-Agent Orchestration**. Proven track record leading cross-functional teams of **15+ engineers**, architecting systems handling **4,500+ daily API requests**, processing **100K+ daily records**, and delivering **18-38% model performance improvements** for **Fortune 500 clients** (Cisco, ITC, KDP, Marco). Deep expertise in end-to-end ML pipelines, containerized deployment, and robust cloud API architecture.

TECHNICAL SKILLS

Python, Django, Flask, FastAPI, Streamlit, LLMs (OpenAI, Anthropic, Vertex AI, Google Gen AI), LangChain, LangGraph, RAG, Multi-Agent Systems, Scikit-Learn, TensorFlow, Keras, PyTorch, Optuna, Transformers, NLP, Apache Airflow, Databases (MySQL, PostgreSQL, MongoDB, Cassandra, Vector DBs), Web Scraping (Playwright, Selenium), Git, Docker, GCP, AWS, CI/CD, MLOps

EXPERIENCE

Senior AI/ML Engineer

R Systems International (Client: Lendistry) | Jun 2026 – Present

- Architecting production Agentic AI workflows and multi-agent coordination using MCP tools and guardrails
- Implementing Harness Engineering and Agentic Ops for automated validation and observability
- Building ML Flow integration and Agentic E2E evaluation frameworks for strategic agent benchmarking

AI/ML Engineer

AI Tech Solutions Ltd. | Nov 2024 – Oct 2025

- Designed and orchestrated multi-agent AI systems utilizing LangGraph for complex task execution
- Enhanced context retrieval relevance by 38% through hybrid RAG implementation
- Built high-throughput, asynchronous backend microservices using FastAPI

Software Engineer

Aimlytics Technology | Sep 2022 – Oct 2024

- Architected scalable RAG-based assistants reliably processing 4,500+ daily API requests
- Fine-tuned BERT and RoBERTa models, improving text classification F1-scores by 18%
- Orchestrated ETL data pipelines via Airflow to handle 100K+ daily records for ML workflows

Machine Learning Engineer Intern

iNeuron.ai | Aug 2021 – Sep 2022

- Accelerated hyperparameter tuning cycles by 35% using Optuna optimization
- Deployed containerized machine learning endpoints using Docker and Flask

PROJECTS

- Omodore**: Autonomous Agentic AI platform demonstrating LangGraph state orchestration
- GALAMBO**: Multi-modal AI assistant integrating vision models and real-time tool use

- **SalesMoji:** Scalable AI sales intelligence system leveraging vector search retrieval

EDUCATION

B.Tech Mechanical Engineering | GATE Qualified